



Artificial Neural Networks

Artificial Neural Networks (ANN)


- ANNs inspired by biological neural systems.
- Human brain consists of nerve cells called neurons
- According to neurologists the **human brain learns by changing the strength of the synaptic connection between neurons upon repeated stimulation by the same impulse.**



General structure of an ANN




Recurrent NN







Feed Forward NN

A Look at a Processing Node

- Sums up weighted inputs into it, adds bias and uses activation function to compute result.




Activation Functions; Examples



Function	Expression
ReLU	$R(z) = \max(0, z)$
Sigmoid	$\sigma(z) = \frac{1}{1 + e^{-z}}$
Hyperbolic tan (tanh)	$\sigma(z) = \frac{e^{-z} - e^z}{e^{-z} + e^z}$
Gaussian/Radial Basis function	$\sigma(z) = e^{-\frac{z^2}{2}}$
Step function	$\sigma(z) = \begin{cases} 0, & z \leq \theta \\ 1, & z > \theta \end{cases}$

Simple Perceptron

- A one-layer feed forward network




Input to activation
function:

$$b + \sum_{i=1}^n w_i x_i$$

Simple Perceptron

- Often, bias term represented as a weight with unit input.




Simple Perceptron

- Used to separate linearly separable patterns.
 - E.g., In 2D, the boundary line is always a straight line...


$$\sum_{i=0}^2 w_i x_i = 0$$

$$\Rightarrow w_1 x_1 + w_2 x_2 + b = 0$$



- Weights determine the slope of the line. Bias determines the y-intercept.

Simple Perceptron



- In 3D (i.e., 3 features, x_1 , x_2 and x_3), discrimination boundary is a 2D plane:
$$w_1x_1 + w_2x_2 + w_3x_3 + b = 0$$
- In general, when we have n features, our decision boundary is $n-1$ dimensional.
- A **hyperplane** is a subspace of one dimension less than its ambient space
- Learning process of a perceptron classifier can be said to amount to learning a hyperplane classifier

Perceptron Learning Algorithm

- Learning Process is basically to find the vector \mathbf{w} of weights \mathbf{w}_0 through \mathbf{w}_n that define a hyperplane separating the classes.
- Algorithm begins with some initial weights vector \mathbf{w}^i
- Algorithm cycles through the training set, a training sample at a time and makes decisions on adjustment of weights with each input sample.
- It's a **mistake-driven** algorithm
 - Only updates \mathbf{w} when it makes a mistake; i.e., when it wrongly predicts the label of the current training example
 - Doesn't update \mathbf{w} when it correctly predicts the label of the current training example

Perceptron Learning Algorithm

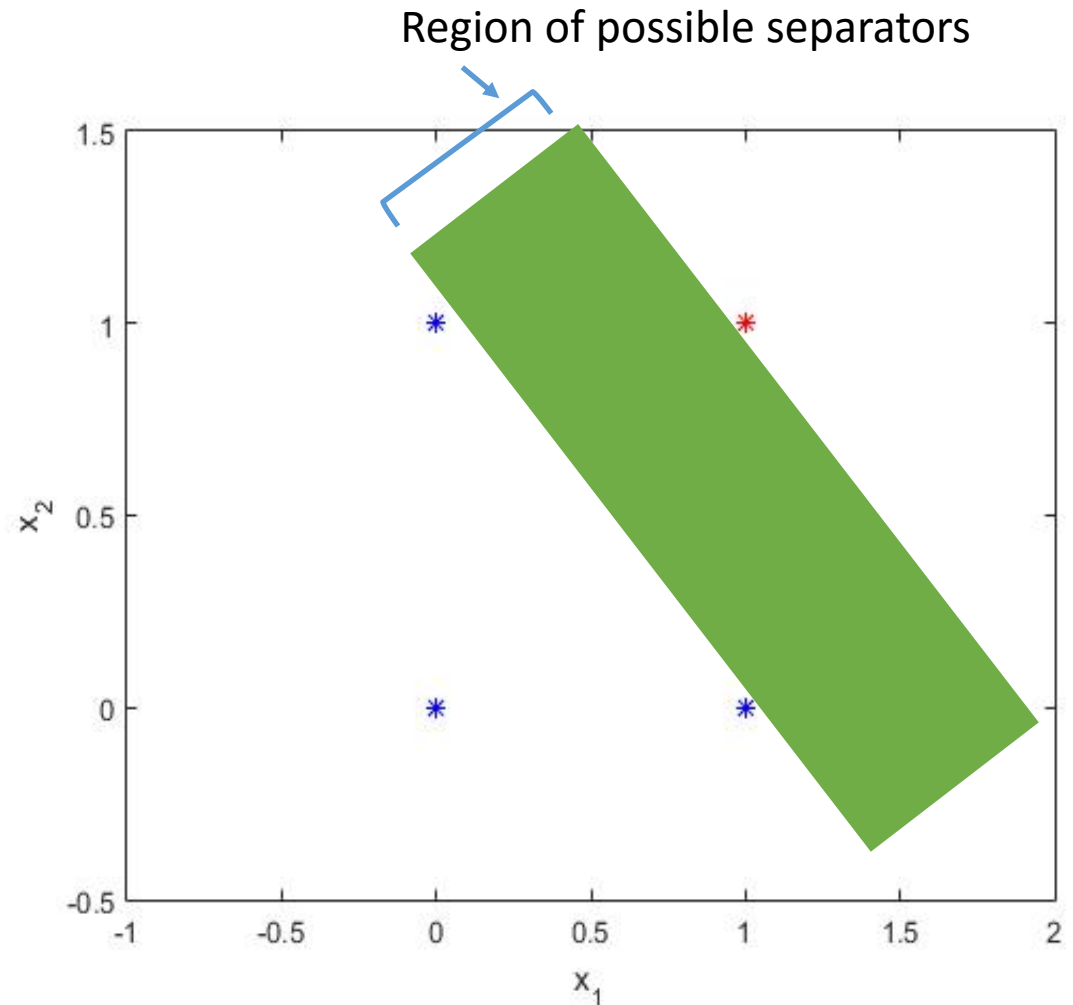
1. Initialize weight vector
2. Select random sample from training set as input
 - ✓ Lets denote each sample as $(x_{\text{train}}, y_{\text{target}})$ where x is the feature vector and y_{target} is the class label of the sample.
3. Compute the output $y_{\text{classifier}}$
4. If $y_{\text{classifier}} = y_{\text{target}}$, do nothing.
5. Otherwise compute the error $\delta = y_{\text{target}} - y_{\text{classifier}}$ and the change in weight $\Delta w = \eta \times \delta \times x_{\text{train}}$ and update the weight vector using $w_{\text{new}} = w_{\text{old}} + \Delta w$
6. Repeat this until the entire training set is classified correctly.

Perceptron Learning Example

- Perceptron that implements a logical **AND**

Input #1	Input #2	Output
0	0	0
0	1	0
1	0	0
1	1	1

- Initial weight vector: $[-2 \ 3 \ 1]$;
 - First component of vector is bias term
- Learning rate: 0.4
- Activation function: $\sigma(z) = \begin{cases} 0, & z \leq 0 \\ 1, & z > 0 \end{cases}$



Perceptron Learning Example

- Notation: $[w_0 \ w_1 \ w_2] = [-2 \ 3 \ 1]$ where w_0 is weight of bias term; Features represented as $[x_0 \ x_1 \ x_2]$, where x_0 is the bias term (recall this equals 1), and $x_1 \ x_2$ are the input features for our training example. Denote, $y_t = y_{\text{target}}$; $y_c = y_{\text{classifier}}$
- Epoch #1

w=[w ₀ w ₁ w ₂]			x=[x ₀ x ₁ x ₂]			y _t	Input, z to Activation function	σ(z)=y _c	δ= y _t -y _c	Δw=[Δw ₀ Δw ₁ Δw ₂]		
-2	3	1	1	0	0	0	-2	0	0	0	0	0
-2	3	1	1	0	1	0	-1	0	0	0	0	0
-2	3	1	1	1	0	0	1	1	-1	-0.4	-0.4	0
-2.4	2.6	1	1	1	1	1	1.2	1	0	0	0	0
-2.4	2.6	1										

Perceptron Learning Example

Epoch #2

$w=[w_0 \ w_1 \ w_2]$			$x=[x_0 \ x_1 \ x_2]$			y_t	z	$\sigma(z)=y_c$	$\delta = y_t - y_c$	$\Delta w=[\Delta w_0 \ \Delta w_1 \ \Delta w_2]$		
-2.4	2.6	1	1	0	0	0	-2.4	0	0	0	0	0
-2.4	2.6	1	1	0	1	0	-1.4	0	0	0	0	0
-2.4	2.6	1	1	1	0	0	0.2	1	-1	-0.4	-0.4	0
-2.8	2.2	1	1	1	1	1	0.4	1	0	0	0	0
-2.8	2.2	1										

Perceptron Learning Example

Epoch #3

$w=[w_0 \ w_1 \ w_2]$			$x=[x_0 \ x_1 \ x_2]$			y_t	z	$\sigma(z)=y_c$	$\delta = y_t - y_c$	$\Delta w=[\Delta w_0 \ \Delta w_1 \ \Delta w_2]$		
-2.8	2.2	1	1	0	0	0	-2.8	0	0	0	0	0
-2.8	2.2	1	1	0	1	0	-1.8	0	0	0	0	0
-2.8	2.2	1	1	1	0	0	-0.6	0	0	0	0	0
-2.8	2.2	1	1	1	1	1	0.4	1	0	0	0	0
-2.8	2.2	1										

Final Weight Vector, $w = [-2.8 \ 2.2 \ 1]$



Perceptron Convergence Theorem

- Given data with linearly separable classes, the perceptron learning rule is guaranteed to find the separating hyperplane in a finite number of iterations
 - Requirement: Learning rate sufficiently small


Under the hood – Why the perceptron rule works

- Preliminaries:
 - Suppose two vectors \mathbf{a} and \mathbf{b} are separated by an angle θ . The inner product $\mathbf{a} \cdot \mathbf{b}$ is defined as $\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos \theta$.
 - L2 norm is always **non-negative**; -- it can either be positive or zero. This implies the sign of $\mathbf{a} \cdot \mathbf{b}$ can be inferred from the sign of $\cos \theta$. If $\cos \theta$ is negative, $\mathbf{a} \cdot \mathbf{b}$ will be negative; if $\cos \theta$ is positive, $\mathbf{a} \cdot \mathbf{b}$ will be positive.


All **S**tudents **T**ake **C**hemistry !!



$$\cos \theta > 0 \Rightarrow \mathbf{a} \cdot \mathbf{b} > 0$$




$$\cos \theta = 0 \Rightarrow \mathbf{a} \cdot \mathbf{b} = 0$$




$$\cos \theta < 0 \Rightarrow \mathbf{a} \cdot \mathbf{b} < 0$$

Under the hood – Why the perceptron rule works

- Some more preliminaries ...
- Consider two 2D vectors for simplicity. Geometrically, adding two vectors is equivalent to appending one vector to the end of the other.





$$w = v + u$$



$$w = v + -u$$

Under the hood – Why the perceptron rule works

- Some more preliminaries ...
- What if we want some guarantees on the result of the addition, e.g., if we want to add some vector to w and have the resultant vector (i) pointing more towards the direction of some vector v , (ii) pointing further away from the direction of some vector v



Under the hood – Why the perceptron rule works

- Recall the input to the activation function was:

$$\sum_{i=0}^n w_i x_i$$



which for the 3D case, for example, would be equal to:

$$w_0 x_0 + w_1 x_1 + w_2 x_2 + w_3 x_3$$

- But this is the dot product between two vectors $w = [w_0 \ w_1 \ w_2 \ w_3]$ and $x = [x_0 \ x_1 \ x_2 \ x_3]$
- So, we basically have $\sum_{i=0}^n w_i x_i = w \cdot v$
- So, during the perceptron learning process, all that is happening is a computation of a dot product between the weight vector and each input vector**

Under the hood – Why the perceptron rule works

- Lets analyze the cases where the algorithm gave us a wrong result
 - Case 1: Learning algorithm assigned a sample to the class label of 1 (i.e., $w.v$ was >0) yet it should have belonged to class label 0 (i.e., $w.v$ should have been <0).
 - In this case, it means that we had the **weight** and **input vectors** aligned as the figure to the left, yet correct classification would have required them to be aligned as the figure to the right




Under the hood – Why perceptron rule works

- Case 1 cont'd

What we have:

What we need for correct classification of sample



How do we fix Case 1?

- ✓ Rotate vector w away from vector v .
- ✓ That is, find a new value, w_{new} of w , such that:
$$w_{\text{new}} = w + (-\alpha v)$$


Check: How did the perceptron algorithm fix Case 1?

- ✓ In our Case 1, $y_{\text{classifier}} = 1$ (because the classifier finds that $w \cdot v > 0$). Since we know that the classifier fails to get the correct decision, this means $y_{\text{target}} = 0$. Recall that the error is $\delta = y_{\text{target}} - y_{\text{classifier}}$. So, that particular iteration would have $\delta = -1$.
- ✓ And the weight update would be $\Delta w = \eta \times \delta \times x_{\text{train}}$ which is the same as $\Delta w = (-1) \times \eta \times x_{\text{train}}$. The new weight is hence $w_{\text{new}} = w + (-\eta \times x_{\text{train}})$ which is in agreement with our strategy above for fixing Case 1, since η is a positive constant between 0 and 1 and our v is x_{train} .


Under the hood – Why perceptron rule works

- Lets analyze the cases where the algorithm gave us a wrong result
 - Case 2: Learning algorithm assigned a sample to the class label of 0 (i.e., $w.v$ was <0) yet it should have belonged to class label 1 (i.e., $w.v$ should have been >0).

What we need for correct classification of sample:



What we have:



This time we want to align w more towards the direction of v so as to work towards attaining the condition **$w.v > 0$**

Using an idea similar to what we used in the previous case, we need to get $w_{\text{new}} = w + (\alpha v)$.
Easy to check that the perceptron algorithm did just this !

Simple Perceptron: some issues

- On test data, the perceptron will use the final weight vector obtained during training.
- But this may not necessarily be a good idea ! One of the “wrong” weight vectors produced before the final “correct” vector could actually perform better on test data
 - **Averaged perceptron** (average some of the well performing weights and predict)
 - **Voted perceptron** (vote on predictions of intermediate vectors)

Simple Perceptron: some issues

- Perceptron rule won't converge if classes are not linearly separable.
- E.g., XOR problem (OR without AND)

x1	x2	y
0	0	0
1	0	1
0	1	1
1	1	0

